

Sofia Lindqvist

Machine Learning Engineer

Kastanienallee 9, 10435 Berlin · +49 30 23125 066 · sofia.lindqvist@example.com · github.com/slindqvist

About

ML engineer with a research background in statistical learning. I take models from notebook to production: data pipelines, training, evaluation and serving. Recent work on retrieval and fine-tuning of language models for document processing.

Experience

Machine Learning Engineer, Northwind Analytics GmbH (Berlin)

03/2023 – present

- Built a document classification service for insurance claims (PyTorch, Hugging Face, FastAPI); 94 % accuracy on 40 classes
- Designed an evaluation harness for LLM-based extraction with gold labels and regression tests
- Orchestrated training and batch scoring with Airflow on Kubernetes; experiment tracking with MLflow

Data Scientist, Fjordline Insurance AB (Stockholm, Sweden)

09/2020 – 02/2023

- Churn and pricing models with gradient boosting (XGBoost, scikit-learn)
- Fraud detection prototype that flagged 2.3× more confirmed cases than the rule system
- Dashboards for business teams in Tableau; SQL on Snowflake

Research Assistant, KTH Royal Institute of Technology (Stockholm)

01/2018 – 08/2020

- Research on uncertainty estimation in neural networks; two workshop papers
- Teaching assistant for "Machine Learning, Advanced Course"

Education

MSc Machine Learning, KTH Royal Institute of Technology, 2016 – 2018

BSc Mathematics, Uppsala University, 2013 – 2016

Skills

Python, SQL, PyTorch, scikit-learn, XGBoost, Hugging Face Transformers, pandas, NumPy, Airflow, MLflow, FastAPI, Docker, Kubernetes, Snowflake, Tableau, LLM fine-tuning, RAG

Languages

Swedish (native) · English (C2) · German (B2)